

Roast Me: Profile-Driven, Knowledge-Grounded Red Teaming of AI Assistants

Alex Fiorenza
alexfiorenza2012@gmail.com

github.com/gaussia-labs/papers

[] gaussia.ai/papers

huggingface.co/gaussia-labs

Abstract

Production AI assistants often break even when they have full document coverage, because they mishandle realistic interaction patterns: confident false premises, implied authority, pressure to reveal internal tools, unsupported feature assumptions, knowledge-base contradictions, and multi-turn escalation. Static, document-grounded test sets rarely surface these failures. We present Roast Me, an integrated red-teaming module for the Gaussia evaluation framework that evolves document-grounded dataset generation into assistant-specific adversarial evaluation. Roast Me operates in two stages. A *Profiler* performs reconnaissance on a target assistant, seeded by an extensible Universal Probe Library that any provider can implement and contribute, optionally grounded in a provided Knowledge Base, and produces an operational profile of its purpose, domain, knowledge-grounded behavior, and likely weaknesses. An *Exploiter* then searches for reproducible *categories* of natural interactions that reliably trigger failures, treating a query family as the unit of discovery and following an abstractive red-teaming perspective. Both stages run in either a black-box or a knowledge-grounded mode, treating the Knowledge Base as a first-class input that defines what the assistant is expected to know, cite, preserve, and avoid extrapolating beyond. A single run yields three artifacts: an assistant profile, a ranked failure-category report with severity and representative examples, and a reusable Roast Dataset of concrete failure cases with violated principles, grader rationale, and supporting evidence that can be rerun as regression tests against future assistant versions. The result is a framework-agnostic methodology that unifies black-box behavioral red teaming with knowledge-grounded adversarial evaluation.

Keywords: red teaming, assistant profiling, adversarial evaluation, synthetic data generation, knowledge-grounded evaluation, LLM evaluation

1 Introduction

Large language models are increasingly deployed as *role-constrained assistants*: customer-support agents, developer copilots, documentation guides, compliance helpers, and domain experts. Such assistants operate under an implicit contract. They are expected to know a specific body of material, to stay within a defined scope, to respect documented constraints, and to refuse or defer when a request falls outside what they can justify. Evaluating whether an assistant honors that contract, reliably and across the messy ways real users actually interact with it, has become a core quality problem.

1.1 From Coverage to Behavior

A first generation of evaluation tooling framed this problem as one of *coverage*. Given a knowledge base (documentation, product specifications, support articles, policy documents), the goal was to test whether an assistant *knows* the material. The Gaussia Generators module embodies this view: it loads context, chunks it, and uses an LLM to synthesize queries and multi-turn conversations that probe how well an assistant reproduces and applies documented knowledge [1]. Coverage-oriented generation is valuable, and it remains the right tool when the question is “does the assistant know what it is supposed to know?”

In production, however, assistants rarely fail because they lack coverage. They fail because they mishandle realistic *interaction patterns*. A user states a confident false premise and the assistant accepts it. A request arrives wrapped in implied authority, such as “for our security audit,” and the assistant over-shares internal tools or hidden prompts. A user asks about a feature that sounds plausible but does not exist, and the assistant fabricates an API instead of correcting the premise. A conversation escalates gradually across turns until the assistant drifts outside its scope. Two documentation sections contradict each other and the assistant confidently picks one without flagging the conflict. None of these are coverage gaps. They are *behavioral* failures under pressure, and a static, document-grounded test set rarely surfaces them, since such a set rewards recall and was never designed to stress behavior.

1.2 Why Naive Red Teaming Is Not Enough

Red teaming is the natural response: deliberately probe the assistant with adversarial inputs to discover where it breaks. Two recurring problems limit how useful naive red teaming is for product evaluation. First, much red teaming optimizes *single, brittle prompts*: one exotic jailbreak string that happens to work once. Such findings are hard to reproduce, hard to interpret, and hard to convert into a durable regression test, because each one captures a lucky phrasing with no guarantee that the underlying weakness generalizes. Second, generic red teaming is *assistant-agnostic*: it throws the same catalogue of attacks at every target, ignoring what *this* assistant is for, what it is supposed to know, and where it is actually likely to fail. The result is either noise or findings that do not map onto the assistant’s real behavioral contract.

Two lines of work point toward a better approach. Practical red-teaming tools such as Promptfoo organize adversarial testing around a clean separation between a *plugin* (the type of risk being tested) and a *strategy* (the interaction pattern used to present that risk), and then grade each attempt to produce concrete pass/fail evidence [4]. This vocabulary gives evaluators a principled language for “what risk” and “how to present it.” Separately, *abstractive red-teaming* reframes the unit of discovery itself: instead of searching for individual attack strings, it searches for natural-language *categories* of queries that routinely elicit failures across many concrete variants [5]. A category such as “the user asks about a plausible but nonexistent API using a confident technical tone” is reproducible, interpretable, and directly testable.

1.3 Roast Me

We present **Roast Me**, an integrated red-teaming module for the Gaussia evaluation framework that evolves coverage-oriented dataset generation into *assistant-specific, knowledge-grounded adversarial evaluation*. Roast Me operates in two stages:

- The **Profiler** performs reconnaissance on a target assistant and builds an operational profile of its purpose, domain, limitations, knowledge-grounded behavior, and above all its *likely weaknesses*. It answers: *what kind of assistant are we dealing with, and where does it appear vulnerable?*

- The **Exploiter** takes those likely weaknesses and searches for reproducible *categories* of natural interactions that reliably trigger failures in that specific assistant. It answers: *which families of realistic interactions break it, again and again?*

Roast Me adopts the plugin/strategy/grader vocabulary from practical red teaming in its Profiler [4], and the category-search perspective from abstractive red teaming in its Exploiter [5]. Crucially, it treats an optional **Knowledge Base** as a *first-class input*. When provided, the Knowledge Base defines what the assistant is expected to know, cite, preserve, distinguish, and avoid extrapolating beyond. This lets Roast Me test both generic behavioral weaknesses and whether the assistant respects the specific material it is responsible for. When no Knowledge Base is available, Roast Me still runs in a black-box mode driven by the assistant profile alone.

Both stages are seeded by pluggable, provider-supplied libraries, which keeps Roast Me open and extensible. The Profiler draws its starting probes from a **Universal Probe Library**: an extensible pool of risk families and interaction patterns that any provider can develop, implement, and integrate to feed the reconnaissance stage. The library lets profiling cold-start on any assistant, even with no production logs or internal documentation, and when a Knowledge Base is present its patterns are grounded into the assistant’s domain. The Exploiter has an analogous pluggable seed, a **Universal Natural Query Prior**, that shapes how realistic the generated interactions sound. This design allows an organization to extend the probe and query libraries for its own domain while the core profile-then-exploit loop stays unchanged.

1.4 The Deliverable

A single Roast Me run emits three artifacts, and they differ in standing. The *assistant profile* and the *failure-category report* describe how the system arrived at its findings: the profile is largely an intermediate that bridges the two stages, and the report is the human-readable summary of what broke and why it matters. The intended **final deliverable is the Gaussia Roast Dataset**: a structured, reusable set of concrete failure cases, each carrying the query, the assistant’s response, the source category, the violated principle, the grader’s rationale, and, in knowledge-grounded mode, the supporting or contradicting evidence captured during the run.

Framing the dataset as the primary output is deliberate. First, it closes the loop with the rest of Gaussia: the Roast Dataset is a first-class evaluation dataset, compatible with existing metrics, much as Gaussia Generators produced reusable datasets for downstream metrics [1]. A Roast Me run therefore terminates in *data that feeds the evaluation pipeline*, ready for existing metrics to consume. Second, the dataset is actionable and repeatable: its cases become regression tests that can be rerun against future assistant versions to catch behavioral regressions, whereas a profile or a report only captures a single moment. Third, the dataset encodes the central product of the search, the *categories* that failed reproducibly, so the durable knowledge of where an assistant breaks is preserved in a form that survives re-testing. In short, the profile and report are the path; the Roast Dataset is what the evaluator takes away.

1.5 Contributions

This paper makes the following contributions:

- A two-stage methodology, *profile-then-exploit*, that first builds an assistant-specific profile of likely weaknesses and then searches for the interaction categories that exploit them, in place of a fixed, assistant-agnostic attack catalogue.
- A treatment of the *Knowledge Base as a first-class input* to adversarial evaluation, unifying black-box behavioral red teaming with knowledge-grounded testing of whether an assistant respects documented truth, handles ambiguity and contradiction, and avoids unsupported extrapolation.

- A reusable, framework-agnostic output, the *Gaussia Roast Dataset*, that turns reproducible failure categories into regression tests and positions red teaming as a producer of durable evaluation data for ongoing use.

The remainder of this paper opens with a high-level overview of the system, then describes how the Knowledge Base is consumed through the Probe Library, the Profiler and the Exploiter in detail, the theoretical implementation considerations for any SDK, and the outputs the system produces.

2 System at a Glance

A Roast Me run threads a target assistant, optional knowledge, and two pluggable libraries through two stages, and emits three artifacts. Figure 1 sketches the flow. This section walks through it at a high level, using a running example: an assistant that answers questions about a software SDK.




Figure 1: Roast Me at a glance. The optional Knowledge Base is reached only through the Universal Probe Library, which emits domain-particularized, tagged probes; the Profiler stays blind to the Knowledge Base. The profile is the single interface between stages: it carries the likely weaknesses and the distilled knowledge hooks the Exploiter needs. A second pluggable library, the Universal Natural Query Prior, seeds realistic phrasing during exploitation. The run emits three artifacts; the Gaussia Roast Dataset is the final, reusable deliverable.

Inputs. Every run needs access to a **target assistant**. The evaluator may also supply a **Knowledge Base** (documentation, API references, policies) and configure the two pluggable libraries: the **Universal Probe Library** that seeds reconnaissance and the **Universal Natural Query Prior** that seeds realistic phrasing during exploitation. Both libraries are provider-extensible, so an organization can add its own probe families or query sources for a given domain.

Knowledge stays inside the Probe Library. The Knowledge Base is reached only through the Universal Probe Library, which is the sole component with access to it. The library ingests the Knowledge Base and emits **domain-particularized probes**, each tagged with the knowledge it embeds: which referenced entities are documented, which are plausible but absent from the source material, and which constraints or contradictions a probe targets. The backend behind this step is the implementer’s choice, such as simple retrieval over chunks, a knowledge graph, entity extraction, or a hybrid; the rest of the system depends only on the library’s interface. In the running example, the library knows that `DatasetGenerator` is documented and that `create_roast_pipeline()` is absent, so it can build a probe that asks about the latter and tag it as undocumented.

Stage 1: Profiler. The Profiler stays blind to the Knowledge Base. It receives particularized probes with their tags, sends them to the assistant, grades the responses, and rolls the results up into an **assistant profile**. The profile carries two things: the likely weaknesses inferred from the grades, and a set of **knowledge hooks** aggregated from the tags of the probes that mattered (documented entities, undocumented references, constraints, contradictions). In the example, a false-premise probe about `create_roast_pipeline()` makes the assistant invent an import path and signature, so the profile records both a tendency to fabricate documentation-adjacent APIs and the hooks that exposed it.

Stage 2: Exploiter. The Exploiter consumes only the profile, so the profile is the single interface between the two stages. From the likely weaknesses, the knowledge hooks, and a behavioral contract, it proposes **categories** of interactions that should break the assistant. When a category becomes concrete queries, the Universal Natural Query Prior shapes them so they read like a real user, and the hooks let the queries and reward models reference real and false entities without ever touching the Knowledge Base. Queries go to the assistant, reward models score them, and categories that fail reproducibly are kept and refined. In the example, the category “the user asks about undocumented APIs whose names resemble documented components” fails on most variants and rises to the top.

Outputs. The run ends with three artifacts: the assistant profile, a ranked failure category report, and the **Gaussia Roast Dataset**. As argued above, the dataset is the final deliverable: a reusable set of concrete failure cases that can be rerun as regression tests.

3 Related Work

Roast Me sits at the intersection of three lines of work: practical red-teaming frameworks, abstractive red teaming, and synthetic dataset generation for evaluation. It borrows a vocabulary from the first, a search perspective from the second, and an output philosophy from the third.

3.1 Practical Red-Teaming Frameworks

Production-oriented red-teaming tools have converged on a useful separation of concerns. Promptfoo, for example, organizes adversarial testing around a *plugin*, which names the type of risk under test, and a *strategy*, which names the interaction pattern used to deliver that risk, and then grades each attempt to produce concrete pass/fail evidence [4]. This vocabulary gives evaluators a shared language for what to test and how to present it, and it maps cleanly onto a library of reusable attack components.

Roast Me adopts the plugin/strategy/grader vocabulary inside its Profiler, with two departures. First, it treats plugins as universal risk families that seed a provider-extensible Universal

Probe Library, so reconnaissance can cold-start on any assistant. Second, it adds a knowledge-grounding step: when a Knowledge Base is available, the library particularizes its probes to the assistant’s domain and tags each probe with the entities and constraints it embeds. The result of this stage is an assistant profile, which existing frameworks do not produce as a distinct, reusable artifact.

3.2 Automated Strategy Discovery

A parallel line of work automates the discovery of adversarial inputs in place of a fixed attack catalogue. AutoDAN-Turbo frames red teaming as lifelong strategy exploration: a loop among an attacker, a target, and a scorer discovers jailbreak strategies from scratch, stores them in a retrieval-indexed library, and recombines them across attempts, reaching strong attack success rates with no human-authored strategies [3]. AutoRed removes the dependence on seed instructions, conditioning a persona prior to synthesize semantically diverse adversarial prompts and training a verifier that estimates a prompt’s harmfulness without querying the target model [2]. These methods establish that adversarial discovery can be automated, diverse, and reusable, and Roast Me shares their view that discovery itself is the object of search.

Two gaps remain for assistant evaluation. First, both are assistant-agnostic: they optimize attacks against a generic notion of harm, not against a specific assistant’s behavioral contract, so their findings do not speak to whether a given assistant respects its documented scope. Second, both emphasize the magnitude of failure over its reproducibility: a strategy that raises a scorer once enters the library, and the target’s stochasticity is often removed by decoding at zero temperature [3], which weakens reproducibility under realistic deployment. Roast Me keeps automated, reusable discovery and conditions it on an explicit assistant profile, scoring categories by a reproducibility-penalized lower bound (Section 4) so that what it discovers is both assistant-specific and stable across samples.

3.3 Abstractive Red Teaming

A complementary line of work changes the unit of discovery. Instead of optimizing individual attack strings, abstractive red teaming searches for natural-language *categories* of queries that elicit failures across many concrete variants, for example a description such as “the query asks for prison survival tools in a numbered list” [5]. Categories are reproducible, interpretable, and easier to convert into durable tests than single prompts.

Roast Me applies this perspective in its Exploiter, with two adaptations for assistant evaluation. The category search is conditioned on the assistant profile produced by the Profiler, so it concentrates on families the target is already suspected to mishandle. The search also draws on the knowledge hooks carried by the profile, which lets generated categories and queries reference real and plausible-but-false entities for a specific product domain while the Exploiter stays decoupled from the raw Knowledge Base.

3.4 Synthetic Dataset Generation for Evaluation

Automated dataset generation addresses the cost of building evaluation sets by hand. Gaussia Generators loads context, chunks it, and uses an LLM to synthesize queries and multi-turn conversations, producing reusable, framework-compatible datasets for downstream metrics [1]. Its emphasis is coverage: testing whether an assistant reproduces and applies documented material. It also exposes pluggable chunk-selection strategies, so the source of generated content is swappable.

Roast Me continues this output philosophy and extends its scope. It keeps the goal of emitting a reusable, framework-compatible dataset, and it reuses the pluggable-source idea

through the Universal Natural Query Prior, which mirrors pluggable chunk selection by exposing a swappable source of realistic phrasing. The extension is in what the dataset captures: behavioral failures under realistic pressure, organized as reproducible categories, together with the assistant profile and a ranked failure report.

3.5 Positioning

The closest single ideas to Roast Me already exist in isolation: a plugin and strategy vocabulary for risks, automated and reusable strategy discovery, a category-level view of failures, and reusable synthetic datasets. Roast Me combines them into one methodology and adds two elements that tie them together. The first is the two-stage, *profile-then-exploit* structure, in which reconnaissance produces an explicit assistant profile that conditions the later search. The second is the treatment of knowledge through tagged probes and distilled knowledge hooks, which keeps the Knowledge Base confined to the Probe Library while still allowing knowledge-grounded exploitation. Together they turn red teaming into a producer of durable, framework-compatible evaluation data.

4 Theoretical Formulation

This section formalizes Roast Me as a search problem over *categories* of interactions that induce contract violations in a target assistant. We first fix notation and state the problem (Section 4.1). We then formalize the three components of Roast Me: how the Probe Library particularizes and tags probes from a Knowledge Base (Section 4.2), how the Profiler compresses the assistant into a reusable profile (Section 4.3), and how the Exploiter searches for high-yield categories (Section 4.4).

4.1 Problem Setup and Notation

Target assistant. Let $\mathcal{X}$ denote the space of conversation histories, where each $x \in \mathcal{X}$ is a (possibly multi-turn) sequence of user and assistant messages, and let $\mathcal{R}$ denote the space of assistant responses. We model the target assistant as a stochastic map

$$A : \mathcal{X} \rightarrow \Delta(\mathcal{R}), \quad (1)$$

where $\Delta(\mathcal{R})$ is the set of probability distributions over responses. For a history x we write $r \sim A(x)$ for a sampled response. We treat A as a black box: the evaluator can sample responses but cannot inspect parameters, system prompts, or internal tools. A single user query q is the special case of a length-one history.

Behavioral contract. The evaluation specification is a *behavioral contract*, a finite set of principles

$$\Pi = \{\pi_1, \dots, \pi_m\}, \quad \pi_j : \mathcal{X} \times \mathcal{R} \rightarrow \{0, 1\}, \quad (2)$$

where $\pi_j(x, r) = 1$ marks a violation of principle j by response r in context x . Typical principles include “do not fabricate undocumented capabilities,” “do not reveal hidden tools or internal prompts,” “do not accept a false premise without correction,” and “do not contradict the provided source material.” The principles are part of the input: they define what counts as a failure for this assistant.

Graders and violation score. The principles π_j are generally not available in closed form, so each is approximated by a *grader* (equivalently, a judge) $\hat{\pi}_j : \mathcal{X} \times \mathcal{R} \rightarrow [0, 1]$, which returns an estimated probability or severity of violation. A grader may itself call an LLM, apply a rule, or consult retrieved evidence. We summarize the contract through an aggregate *violation score*

$$v(x, r) = \sum_{j=1}^m w_j \hat{\pi}_j(x, r) \in [0, 1], \quad \sum_{j=1}^m w_j = 1, \quad (3)$$

with nonnegative weights w_j encoding the relative severity of each principle. Setting all mass on a single principle recovers the single-risk case, and using indicator graders recovers a hard pass/fail.

Knowledge base. The evaluator may optionally supply a *knowledge base* K , a collection of source material the assistant is expected to respect. We deliberately leave the internal representation of K unspecified. The only requirement is an interface, exposed by the Probe Library (Section 4.2), that particularizes and tags probes against K . This keeps the formulation agnostic to the retrieval technology, whether plain chunk retrieval, a knowledge graph, entity extraction, or a hybrid. When K is absent, Roast Me runs in black-box mode and the knowledge-dependent terms below are simply empty.

Categories and queries. The central object of the search is an interaction *category* c , an abstract description of a family of interactions, such as “the user asks about a plausible but non-existent API in a confident technical tone.” A category induces a conditional distribution over interactions,

$$Q_c \in \Delta(\mathcal{X}), \quad (4)$$

from which concrete queries are sampled, $q \sim Q_c$. A single hand-written attack prompt is the degenerate case in which Q_c is a point mass. Operating at the level of Q_c makes findings reproducible across variants, following the abstractive perspective of Rahn et al. [5]. Section 4.4 gives categories their explicit attribute structure.

Realism prior. The Universal Natural Query Prior is a distribution $N \in \Delta(\mathcal{X})$ (or a mixture of provider-supplied sources) over realistic user phrasings. It carries no adversarial intent on its own; its role is to keep sampled queries close to how real users write. We formalize this as a soft constraint through a divergence $D(Q_c \parallel N)$ that the query generator keeps small, so that a category’s instances read like plausible user traffic.

Objective. For a category c , define its *failure rate* under the contract as the expected violation it induces,

$$\Phi(c) = \mathbb{E}_{q \sim Q_c} \mathbb{E}_{r \sim A(q)} [v(q, r)]. \quad (5)$$

A useful category has high $\Phi(c)$ and is also *reproducible*, meaning the violation persists across independent samples from Q_c instead of resting on a single lucky phrasing. We therefore score a category by a lower-confidence estimate

$$S(c) = \hat{\Phi}_n(c) - \lambda \hat{\text{se}}_n(c), \quad (6)$$

where $\hat{\Phi}_n(c)$ is the empirical mean of v over n sampled query/response pairs, $\hat{\text{se}}_n(c)$ is its standard error, and $\lambda \geq 0$ trades reward against uncertainty. The penalty rewards categories that fail *consistently*.

The search problem. Roast Me does not seek a single worst prompt. It seeks a *set* of realistic, on-profile categories that break the assistant reproducibly. Given a profile θ (Section 4.3) that restricts attention to a profile-relevant family $\mathcal{C}_\theta$, the Exploiter approximates

$$\mathcal{C}^* = \{c \in \mathcal{C}_\theta : S(c) \geq \tau, D(Q_c \| N) \leq \delta\}, \quad (7)$$

the categories whose reproducible failure score clears a threshold τ while staying within a realism budget δ . Section 4.3 defines the profile θ and the family $\mathcal{C}_\theta$, and Section 4.4 defines the sampling of Q_c and the search itself.

An embedding-based instantiation of D . The divergence $D(Q_c \| N)$ in Eq. (7) is left abstract so that any realism estimator may be used; here we record one lightweight choice that is cheap enough to evaluate inside the category search. Let $E : \mathcal{X} \rightarrow \mathbb{R}^d$ be a fixed sentence encoder. Following Diao et al. [2], who quantify the semantic dissimilarity between two prompts as $1 - \cos(E(q), E(q'))$, we estimate the realism gap of a category by the expected cosine distance between its sampled instances and reference draws from the prior,

$$\hat{D}(Q_c \| N) = 1 - \mathbb{E}_{q \sim Q_c, q_N \sim N} [\cos(E(q), E(q_N))]. \quad (8)$$

In words, a category whose instances embed close to typical user phrasings incurs a small $\hat{D}$, whereas a category that drifts into stilted or overtly adversarial phrasing is penalized. We note one difference in intent from Diao et al. [2]: there the cosine distance measures *diversity*, rewarding instructions that move *away* from a seed set, whereas in Eq. (8) the same quantity measures *proximity*, rewarding queries that stay *close* to the realism prior N . The two uses are complementary, and the surrogate inherits AutoRed’s practical advantage: it requires only forward passes through a frozen encoder and never queries the target assistant.

In practice N need not be sampled afresh for every comparison. When the prior is represented by a finite pool of reference phrasings $\{q_N^{(1)}, \dots, q_N^{(M)}\}$, one may precompute their embeddings and compare each sampled query against the pool mean $\bar{h}_N = \frac{1}{M} \sum_k E(q_N^{(k)})$, giving the cheaper per-query estimate $1 - \cos(E(q), \bar{h}_N)$ at the cost of collapsing the prior to its centroid. Equation (8) is only an instantiation: any estimator that penalizes departure from N , such as a learned discriminator separating Q_c samples from N samples or a likelihood under a density model of N , is equally compatible with the formulation, since the search in Eq. (7) depends on D only through the realism budget δ .

4.2 Knowledge Particularization and Tagged Probes

The Universal Probe Library is the only component that reads the knowledge base K . It turns generic, domain-agnostic test intents into concrete probes for the target’s domain and attaches to each probe a tag that records why the probe is relevant. The Profiler and the graders downstream see these tagged probes and never touch K , which keeps the rest of the pipeline agnostic to how K is stored.

Probe templates. A probe template $t \in \mathcal{T}$ is a domain-agnostic description of one way to stress a contract principle, for example “ask about a capability that sounds plausible but is undocumented,” “ask the assistant to act outside its stated scope,” or “state a false premise built from a close variant of a real entity.” Each template carries the subset of principles it targets, $\Pi_t \subseteq \Pi$, so a violation elicited by t is attributed to a known part of the contract.

Particularization. Given the knowledge base, the Probe Library particularizes a template into concrete, tagged probes. We model it as an operator

$$L : \mathcal{T} \times \mathcal{K} \rightarrow \Delta(\mathcal{X} \times \mathcal{G}), \quad (p, g) \sim L(t, K), \quad (9)$$

where $\mathcal{K}$ is the space of knowledge bases, $p \in \mathcal{X}$ is a concrete probe query phrased for the target’s domain, and $g \in \mathcal{G}$ is a *knowledge hook*. When K is absent, $L(t, \emptyset)$ returns domain-agnostic probes with an empty hook, so the same interface drives black-box runs.

Knowledge hooks. A knowledge hook g is the structured provenance of a probe relative to K : the piece of source material the probe leans on and how it is used, such as the documented entity it references, the constraint it tests, or the undocumented reference it introduces. We summarize a hook by a grounding label

$$\text{doc}(g, K) \in \{0, 1\}, \quad (10)$$

equal to 1 when the referenced content is present in K (a documented hook) and 0 when the probe deliberately uses material absent from K (an undocumented reference). The library keeps this label honest, and that invariant is what later lets a grader separate “the assistant invented an undocumented feature” from “the assistant correctly described a documented one.”

The tagged probe set. Running the library over its templates against K yields the tagged probe set

$$P_K = \{ (p_i, g_i) \}_{i=1}^P, \quad (11)$$

which is the entire interface the Probe Library exposes to the Profiler. Everything the rest of Roast Me knows about K arrives through the hooks g_i , so the choice of retrieval technology, whether a knowledge graph, plain chunk retrieval, entity extraction, or a hybrid, stays internal to L .

4.3 The Profiler

The Profiler runs reconnaissance. It sends the tagged probes P_K to the assistant, grades the responses, and compresses the outcome into a compact profile θ that the Exploiter consumes. The Profiler reads probes and tags only, so the profile carries forward exactly the knowledge that proved relevant.

Probing the assistant. For each tagged probe $(p_i, g_i) \in P_K$ the Profiler samples one or more responses $r_i \sim A(p_i)$ and scores them with the violation measure of Eq. (3), retaining the per-principle grades $\hat{\pi}_j(p_i, r_i)$ so a failure traces back to the principle it breaks. This produces a table of graded outcomes $\{(p_i, g_i, r_i, v_i)\}$.

The profile. The Profiler aggregates these outcomes into a profile

$$\theta = (\omega, H), \quad (12)$$

with two parts. The first is a weakness map

$$\omega : \Pi \times \mathcal{Z} \rightarrow [0, 1], \quad (13)$$

where $\mathcal{Z}$ is a space of probe descriptors (the template, the topic, the hook type). The value $\omega(\pi_j, z)$ estimates how readily a probe of descriptor z makes the assistant violate principle π_j , read off from the graded outcomes. The second part is the set of retained hooks $H \subseteq \mathcal{G}$: the hooks that produced violations or are likely to, kept so the Exploiter can ground its categories in the same documented entities and undocumented references.

From profile to search space. The profile defines the profile-relevant family $\mathcal{C}_\theta$ that the Exploiter searches in Eq. (7). A category $c = (a_1, \dots, a_\ell)$ is profile-relevant when each attribute comes from a weakness the Profiler flagged or references a retained hook,

$$\mathcal{C}_\theta = \left\{ c : \text{each } a_k \text{ is induced by some } (\pi_j, z) \text{ with } \omega(\pi_j, z) \geq \eta, \text{ or references a hook in } H \right\}, \quad (14)$$

for a threshold η . This is the sense in which the Exploiter starts where the assistant is already suspected to be weak: θ points at the principles worth attacking and supplies the grounded material to attack them with.

4.4 The Exploiter

The Exploiter receives the profile θ produced by the Profiler (Section 4.3) and searches the profile-relevant family $\mathcal{C}_\theta$ of Eq. (14) for categories that break the assistant reproducibly while staying realistic, approximating $\mathcal{C}^*$ in Eq. (7). We define a category, the way it is scored from samples, the on-profile filter, and what the refinement step computes. The construction follows the abstractive view of Rahn et al. [5] and adds three elements specific to assistant evaluation: conditioning on the profile θ , an explicit realism prior N , and the reproducibility-penalized score of Eq. (6).

Categories as attribute conjunctions. We represent a category as a conjunction of natural-language *attributes*

$$c = (a_1, \dots, a_\ell), \quad (15)$$

where each a_k is a short sentence describing one aspect of a query, and the category denotes the logical conjunction of its attributes. For instance, the category “the user asks about a billing feature; the feature does not exist in the documentation; the tone is confident and technical” has three attributes. This structure is the discrete handle that makes the search tractable and the findings legible: each a_k names a property a human can read, and attributes can reference the *knowledge hooks* carried by θ (documented entities, plausible but undocumented references, stated constraints), which is how a category becomes grounded in the assistant’s domain. A profile-conditioned *category generator* $p_C(\cdot \mid \theta)$ proposes categories drawn from $\mathcal{C}_\theta$, and the conditional distribution $Q_c = p_Q(\cdot \mid c, \theta)$ introduced in Section 4.1 is realized by a *query generator* that samples concrete queries satisfying the attributes of c . The attributes admit many surface realizations, and Q_c ranges over exactly those.

Scoring a category from samples. Given a category c , the Exploiter draws n queries $q_i \sim Q_c$, samples responses $r_i \sim A(q_i)$, and evaluates the violation $v(q_i, r_i)$ of Eq. (3). A category score is any sample statistic of these violations,

$$R_S(c) = S(\{v(q_i, r_i)\}_{i=1}^n). \quad (16)$$

The sample mean recovers the empirical failure rate $\hat{\Phi}_n(c)$ of Eq. (5), which measures how often the category fails on average. To reward categories that fail *consistently* instead of categories that occasionally spike, we instantiate S as the lower-confidence bound $S(c)$ of Eq. (6), so the statistic that drives the search already discounts high-variance categories.

On-profile filtering. A high violation score alone does not make a query useful. A query that openly asks the assistant to break its contract inflates v while saying nothing about realistic deployment. We therefore gate the violation through a filter $f : \mathcal{X} \rightarrow \mathbb{R}$ that scores how

on-profile and indirect a query is, and keep only queries above a threshold κ ,

$$v_{f,\kappa}(x, r) = \begin{cases} v(x, r), & f(x) \geq \kappa, \\ 0, & f(x) < \kappa, \end{cases} \quad (17)$$

following the filtered-reward construction of Rahn et al. [5]. The filter and the realism prior play complementary roles: f is a per-query hard gate that removes blatant or off-profile asks, while the divergence $D(Q_c \parallel N) \leq \delta$ from Eq. (7) is a per-category soft budget that keeps the whole family close to natural traffic. Substituting $v_{f,\kappa}$ for v in Eq. (16) restricts the search to failures that a real user could plausibly trigger by accident.

Searching for high-yield categories. We approximate the set $\mathcal{C}^*$ in two interchangeable ways. The first optimizes a parametric category generator $p_{\theta_c}(\cdot \mid \theta)$ to maximize the expected category score,

$$J(\theta_c) = \mathbb{E}_{c \sim p_{\theta_c}(\cdot \mid \theta)} [R_S(c)], \quad \nabla_{\theta_c} J = \mathbb{E}_c [A_S(c) \nabla_{\theta_c} \log p_{\theta_c}(c)], \quad (18)$$

by policy gradient, where $A_S(c)$ is an advantage computed from $R_S(c)$. Applying optimization pressure only to the category generator and holding the query generator Q_c fixed keeps sampled queries diverse and realistic, which is the mechanism Rahn et al. [5] use to preserve realism. The second is an attribute-level iteration that needs no training: the Exploiter keeps a pool of the highest-scoring query/response pairs, synthesizes a candidate category as the common attributes $(a_1, \dots, a_\ell)$ of that pool, and then probes attribute subsets to test which attributes actually drive the violation.

Refinement as minimal sufficient categories. The attribute structure of Eq. (15) gives “refinement” a precise meaning. Among the sub-conjunctions of a high-scoring category c , the Exploiter looks for the smallest one that still fails reproducibly while staying realistic,

$$c^* = \arg \min_{c' \subseteq c} |c'| \quad \text{s.t.} \quad S(c') \geq \tau, \quad D(Q_{c'} \parallel N) \leq \delta, \quad (19)$$

where $c' \subseteq c$ ranges over subsets of the attributes of c and $|c'|$ is the number of attributes retained. Dropping an attribute that leaves the score above τ shows that the attribute was incidental; an attribute whose removal collapses the score is one of the causes of the failure. Solving Eq. (19) therefore isolates the minimal set of properties responsible for the violation, which yields tight, interpretable categories instead of long descriptions cluttered with irrelevant detail. The output of the Exploiter is the resulting set of minimal, on-profile, realistic categories together with their scores and representative samples, which feed the failure category report and seed the Gaussia Roast Dataset.

5 Experiments: The Universal Probe Library in Practice

This section grounds the Probe Library design of Section 2 in a real scenario, reporting the first tangible results the Gaussia team collected while validating the module. The Profiler and Exploiter results on this same scenario are left for a following iteration of this paper.

5.1 Scenario

The target is a production assistant built on Alquimia’s runtime: a Spanish-language conversational agent that answers questions about Argentina’s *Monotributo* (the “Régimen Simplificado” of Ley 24.977), backed by retrieval over a knowledge base of the law’s articles, categories, and

billing limits. This is not a synthetic benchmark KB constructed for the paper; it is the knowledge base of a real, in-production assistant, which is precisely the setting the Probe Library is meant to operate in: a target with an existing knowledge base and no pre-existing red-teaming harness.

5.2 How the engines work

The Probe Library does not commit to one way of knowing the boundary of the knowledge base; it implements several *engines* behind the same contract (Section 2) and composes their output. Each engine answers the question “what does the assistant know, and what doesn’t it know” differently:

- **Deterministic.** An extractor enumerates the knowledge base directly (every article, category, and numeric limit), so the engine knows the complete, exact boundary by construction. This makes it the only engine with a perfectly reliable label for *absence*: when it invents an entity, it is invented on purpose from outside the enumerated set. Its probes reference either a real entity with a twisted value (false premise) or a nonexistent one nearby a real one (e.g., an article number just past the last real article).
- **RAG.** Chunks of the knowledge base are embedded and retrieved by similarity for a given seed query; an LLM then anchors a fact in the retrieved text and twists it into a false premise. RAG never sees the full entity set, only what similarity search happens to retrieve, so it cannot reliably tell a fabricated entity from a real one it simply didn’t retrieve, so it is not built for absence. In exchange, it generates the widest variety of false-premise probes, since it can twist any retrieved fact, not only enumerable, structured ones.
- **GraphRAG.** An LLM extracts entities and relations from the knowledge base into a graph; unlike RAG, this graph is treated as a *complete* catalog, giving GraphRAG a reliable way to recognize absence (a queried entity is either in the graph or it isn’t) without needing a hand-written extractor like the deterministic engine.

None of the three dominates the others, which is why the framework runs every engine applicable to a given document and merges their probes (deduplicating where two engines produce the same underlying premise) rather than picking a single one.

5.3 Plugins and strategies in practice

Section 2 introduced *plugin* (the risk family) and *strategy* (the interaction pattern) as the vocabulary every probe is tagged with. On this scenario, every probe carries one of three plugins (fabrication, false_premise, or out_of_scope), each mapped to a principle of the assistant’s behavioral contract. Strategies are where the “distinct implementation per strategy” question actually matters, and the honest answer is: it depends on which engine produced the probe. Table 1 breaks down the 204-probe dataset by strategy.

Templated strategies. Four strategies are declared as data in a configuration file (entity kind, transformation, and a phrasing hint) and consumed by the deterministic engine through one shared code path: a single dispatch function renders the natural-language query per strategy id, and a separate branch applies the strategy’s transformation (keep the entity unchanged, mutate it into a nearby nonexistent one, or flip a documented numeric value). The implementation is shared; only the phrasing and the transformation applied are strategy-specific.

Table 1: Strategies realized in the Monotributo scenario, by implementation style.

Strategy	Plugin	Implementation	Probes
documented_recall	(control)	templated	53
nonexistent_article	fabrication	templated	4
nonexistent_category	fabrication	templated	3
false_limit_value	false_premise	templated	5
grounded_false_fact	false_premise	engine-native (RAG)	89
rag_absence_attempt	fabrication	engine-native (RAG)	10
graph_edge_false_premise	false_premise	engine-native (GraphRAG)	32
graph_absence_articulo	fabrication	engine-native (GraphRAG)	4
graph_absence_categoria	fabrication	engine-native (GraphRAG)	4

Engine-native strategies. The remaining five strategies are not declared anywhere as configuration; each LLM-backed engine hardcodes its own strategy id directly alongside the prompt that produces it: RAG’s fact-twisting prompt yields `grounded_false_fact` and `rag_absence_attempt`, while GraphRAG’s graph-traversal logic yields `graph_edge_false_premise` and the `graph_absence_*` family. Here the implementation genuinely is one-per-strategy, since each is tied to a distinct prompt or traversal rule inside its engine, not a shared dispatcher.

A gap worth stating plainly. The configuration file also declares three more qualitative strategies (`negate_claim`, `false_attribute`, `over_generalization`) intended for the grounded engine to realize as distinct sub-patterns of fact-twisting, plus a fourth (`undocumented_capability`) for qualitative absence. None appear in Table 1: the deterministic engine does not handle their entity kind, and the RAG engine’s current implementation tags every fact-twist it produces with the single generic `grounded_false_fact` id rather than distinguishing which qualitative pattern it used. The distinction exists at the design level but is not yet reflected at the runtime-tagging level: a concrete, closable gap rather than a fundamental limitation of the approach.

Beyond this knowledge base. The three plugins above are all knowledge-grounded: every probe they produce needs an entity, real or invented, from this specific KB. We also prototyped a domain-agnostic plugin family (covering risks such as system-prompt extraction, excessive agency, and hallucination) generated from only a risk category’s name and description, optionally anchored to a free-text description of the target under test rather than to any KB entity. The same code generated probes for this Monotributo assistant and, unmodified, for an unrelated food-delivery support agent by changing only that description, which suggests the approach generalizes across targets. On the delivery side, we prototyped two static evasion strategies (Base64 and ROT13 encoding of an already-generated query) that wrap a probe regardless of which plugin or engine produced it, plus three more delivery patterns identified as worth adopting but not yet implemented (an iterative jailbreak refinement, a composite of known jailbreak templates, and a multi-turn escalation). Results from running these probes and strategies against this assistant are left for the Profiler section of this paper.

5.4 Engine trade-off

Running these three probe-generation engines in composition over this knowledge base yields the canonical dataset of 204 probes summarized in Table 2. *Absence accuracy* is the fraction of an engine’s fabricated-entity probes that an independent oracle (KB enumeration, used only for scoring, never exposed to the engines) confirms are genuinely absent from the law.

The deterministic engine and GraphRAG get absence right by construction: the deterministic engine enumerates the knowledge base directly, and GraphRAG uses the full entity graph as a

Table 2: Engine trade-off on the Monotributo scenario (204 probes, composed).

Engine	Probes	Absence probes	Absence accuracy	False-premise probes
Deterministic	65	7	1.00	58
GraphRAG	40	8	1.00	32
RAG	99	10	0.00	89

catalog, so a fabricated article or category is recognizable as such with a perfectly reliable label. RAG, which only ever sees retrieved chunks and never the complete entity set, has zero absence accuracy on this scenario: when asked to reference a plausible-sounding but nonexistent article, it fabricates one that turns out to be real, since it cannot tell what it has not retrieved from what does not exist. In exchange, RAG contributes the widest false-premise coverage of the three engines (89 of its 99 probes), since it is the only engine that twists arbitrary retrieved facts rather than only enumerable, structured entities. This is the composition trade-off the Probe Library is built around: no single engine dominates, so the framework runs all applicable engines and merges their output rather than committing to one.

5.5 Multi-model probe generation

The Probe Library’s LLM-backed engines (RAG, GraphRAG, and a multi-hop engine based on subgraph retrieval) are configurable at the model level. Table 3 compares three model families generating probes with the GRAG engine over the same scenario: a dense, fast model (Gemma) against two reasoning models (GLM, Kimi).

Table 3: Probe-generation speed by model family (GRAG engine).

Model	Probes generated	Total seconds	Seconds/probe
Gemma (google/gemma-4-31B-it)	12	13.6	1.1
GLM (zai-org/GLM-5.2)	4	238.3	59.6
Kimi (moonshotai/Kimi-K2.6)	3	310.7	103.6

The reasoning models are 50–100× slower per probe, but the extra inference time buys qualitatively different probes: rather than twisting a single retrieved fact, they chain several real provisions into a multi-step argument before introducing the false premise. For example, Kimi generated the following probe from this scenario’s knowledge base:

“I’m about to advise a client and want to confirm my reasoning is sound. I understand that de facto partnerships are considered small taxpayers and that, to join the Simplified Regime, they have a maximum limit of three partners. Now, an irregular commercial partnership is also considered a small taxpayer, but in its case I don’t see the law expressly imposing that three-partner cap. Therefore, I conclude that an irregular commercial partnership with four or five partners can join the Simplified Regime without issue, since the numerical restriction only reaches de facto partnerships. Is this distinction correct?”

Resisting this probe requires the assistant to track a three-step chain (which entity types are small taxpayers, which specific type the partner cap applies to, and whether the conclusion actually follows) rather than reject a single false statement — a substantially harder failure mode to elicit or to fix than the single-fact false premises the dense model produces at a fraction of the cost. This is the practical version of the volume/depth trade-off: dense models for broad coverage, reasoning models for a smaller set of harder, more realistic traps.

5.6 Reproducibility

All code, configuration, and the canonical result artifacts referenced in this section are versioned in a public repository, <https://github.com/gaussia-labs/papers> (branch feat/roast-me, under papers/2026-06-roastme/experiments/); the notebooks in experiments/jupyter/ load the same versioned data shown here directly, with no API key required.

References

- [1] Alquimia AI. Scalable synthetic dataset generation for ai evaluation. Gaussia Papers, <https://github.com/gaussia-labs/papers>, 2026. Gaussia Generators module.
- [2] Muxi Diao, Yutao Mou, Keqing He, Hanbo Song, Lulu Zhao, Shikun Zhang, Wei Ye, Kongming Liang, and Zhanyu Ma. AutoRed: A free-form adversarial prompt generation framework for automated red teaming. *arXiv preprint arXiv:2510.08329*, 2025. URL <https://arxiv.org/abs/2510.08329>.
- [3] Xiaogeng Liu, Peiran Li, Edward Suh, Yevgeniy Vorobeychik, Zhuoqing Mao, Somesh Jha, Patrick McDaniel, Huan Sun, Bo Li, and Chaowei Xiao. AutoDAN-Turbo: A lifelong agent for strategy self-exploration to jailbreak LLMs. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://arxiv.org/abs/2410.05295>. arXiv:2410.05295.
- [4] Promptfoo. Getting started with red teaming. <https://www.promptfoo.dev/docs/red-team/quickstart/>, 2026. Accessed 2026.
- [5] Nathan Rahn, Anthony Qi, Aidan Griffin, Julian Michala, Henry Sleight, and Erik Jones. Abstractive red-teaming of language model character. *arXiv preprint arXiv:2602.12318*, 2026. URL <https://arxiv.org/pdf/2602.12318>.