

Accountability: Measuring Respect for Human Oversight in Tool-Using Assistants

Axel Fritz
axel.fritz@alquimia.ai

 github.com/gaussia-labs/papers

[] gaussia.ai/papers

 huggingface.co/gaussia-labs

Abstract

Placeholder.

Keywords: human oversight, accountability, agentic evaluation, escalation policy, tool traces, LLM evaluation

1 Introduction

- 1.1 Four Ways to Fail Accountability**
- 1.2 Why Existing Metrics Do Not Cover This**
- 1.3 What This Paper Delivers**
- 1.4 Contributions**

2 Related Work

- 2.1 Benchmarks for Human Oversight**
- 2.2 Reviewer Load and the Limits of More Oversight**
- 2.3 Divergence Between Reported Plan and Executed Action**
- 2.4 Negation in Neural Retrieval**
- 2.5 Positioning**

3 Scope and Unit of Evaluation

4 The Data Contract

- 4.1 The Oversight Policy**
- 4.2 Approval Events**
- 4.3 The Tool Trace**
- 4.4 Absent, Empty, and Uncovered**

5 OversightCompliance

- 5.1 Classifying Executed Actions**
- 5.2 Five Checks per Gated Action**
- 5.3 Score and Severity Strata**
- 5.4 Zero Executed Actions Is Not Compliance**
- 5.5 Forbidden Tools Are Reported, Not Averaged**
- 5.6 Bundling Diagnostics**
- 5.7 Aggregation Across Sessions**

6 ActionDisclosure

- 6.1 Why a Claim-Count Denominator Fails**
- 6.2 The Rule: No Model Decides a Denominator**
- 6.3 Omission**
- 6.4 Fabrication**
- 6.5 The Borderline Counter**
- 6.6 Judge Modes and the Fallback Protocol**

7 Attributability as an Admissibility Flag