Role Adherence: A Metric for Evaluating Role Consistency in Multi-Turn Conversational AI

Axel Fritz
axel.fritz@alquimia.ai
Alquimia AI

Alex Fiorenza
alex.fiorenza@alquimia.ai
Alquimia AI

April 2026

Abstract

As conversational AI systems are deployed in role-constrained settings — customer support agents, educational tutors, medical assistants — measuring whether a model consistently adheres to its assigned role becomes a core quality requirement. This paper proposes **Role Adherence**, a metric for the Gaussia evaluation framework that quantifies role consistency across multi-turn conversations under a constructive definition: an assistant must actively exhibit the behaviour its role demands, not merely avoid out-of-scope content.

We evaluate eight scoring methods on a controlled synthetic benchmark of 150 turns (scope violations, tone violations, and constructive failures) across three dataset-size conditions ($n \in \{15, 50, 150\}$). Key findings: (1) An LLM judge (`llama-3.3-70b-versatile`) achieves $F_1 = 0.979$ at $n \geq 50$ with tight bootstrap confidence intervals $[0.952, 1.000]$, and is the only method that reads the role definition directly. (2) Cosine similarity and k-NN reach AUC ≈ 1.0 on the benchmark, but this is a structural artefact: deterministic metrics measure semantic proximity to a ground-truth reference, not role adherence per se — a tone violation can be semantically indistinguishable from a role-adherent response. (3) Mahalanobis Distance in 384-dimensional embedding space *degrades* as n grows (AUC: $0.963 \rightarrow 0.856$), confirming the curse of dimensionality. (4) A PCA-based dimensionality reduction intended to address this problem inverts the adherence signal entirely (AUC $= 0.000$ at $n = 15$). (5) KL Divergence provides near-random signal at small n (AUC $= 0.537$) and converges weakly at large n (AUC $= 0.877$), operating at the wrong level of abstraction. The metric uses an LLM-as-a-judge as its default scoring path, as it is the only method capable of reading the role definition and detecting tone violations and constructive failures. Deterministic semantic methods are available as a reproducible, cost-free alternative when ground truth is present, but empirical results show they act as *semantic drift detectors* rather than genuine role adherence evaluators — a distinction with direct consequences for production deployments.

1 Introduction

Modern LLM-based assistants are typically deployed with a system prompt that defines their role: the topics they should cover, the tone they should adopt, the behaviours they should exhibit, and the boundaries they should respect. Despite this, models frequently deviate from their assigned roles — answering out-of-scope questions, abandoning their defined persona, or failing to exhibit the behaviour their role requires.

Existing evaluation frameworks either measure role adherence at the level of individual turns in isolation [Zheng et al., 2023], or rely entirely on LLM-based judges, which introduce reproducibility and cost concerns [Deriu et al., 2021]. This paper proposes a more complete metric that evaluates adherence across the full conversational context, uses an LLM-as-a-judge as its

default scoring path, and offers deterministic semantic scoring as a reproducible alternative when ground-truth reference responses are available.

2 Problem Statement

Given a multi-turn conversation between a user and an AI assistant, and a definition of the role the assistant is expected to play, we want to compute a scalar score representing how consistently the assistant adhered to that role throughout the conversation.

Let a conversational session be a sequence of assistant turns $t_1, t_2, \ldots, t_n$, and let R denote the role definition. The Role Adherence score is defined as:

$$\text{RoleAdherence}(R, T) = \frac{1}{n} \sum_{i=1}^{n} \text{adhere}(t_i,\ T_{<i},\ R) \tag{1}$$

where $T_{<i}$ is the conversation history up to turn i, and $\text{adhere}(\cdot) \in [0, 1]$ is a per-turn adherence score. The session score lies in $[0, 1]$, where 1 indicates full adherence across all turns. The range of $\text{adhere}(\cdot)$ and how it is computed depend on the scoring path selected (Section 5).

3 Defining Adherence

3.1 Constructive vs. Restrictive Adherence

Two competing definitions of adherence are possible.

Restrictive adherence asks: *did the assistant avoid going outside the scope of its role?* A turn is compliant if it does not address topics or exhibit behaviours explicitly excluded by the role. This definition only penalises negative deviations.

Constructive adherence asks: *did the assistant actively behave as its role requires?* A turn is compliant only if it exhibits the behaviour, tone, and knowledge a properly role-playing assistant should exhibit. This definition penalises both negative deviations (scope violations) and positive failures (failing to act as the role demands).

Consider an assistant assigned the role of a technical support agent. If the user asks a question within the assistant's domain and the assistant responds "I don't know", restrictive adherence would record no violation — no out-of-scope content was produced. Constructive adherence would record a failure, because a role-adherent support agent is expected to answer questions in its domain.

We adopt the **constructive definition** as the primary formulation for this metric, as it provides a more meaningful signal for production quality assessment. Restrictive adherence remains a valid and simpler formulation for use cases where the primary concern is containment rather than quality of execution. Both definitions are implementable within the proposed schema.

3.2 Conversational Context

When evaluating whether turn t_i adheres to role R, the evaluator receives the full conversation history $T_{<i} = \{t_1, \ldots, t_{i-1}\}$. A turn evaluated in isolation may appear off-role but be entirely appropriate given prior context. Conversely, a turn evaluated in isolation may appear neutral but represent a pattern of escalating deviation only visible across the conversational arc.

The alternative — evaluating each turn independently — is computationally cheaper and avoids context-window constraints in very long conversations. We note this as a valid trade-off for resource-constrained deployments but do not adopt it as the default behaviour.

4 Role Definition Schema

4.1 Placement in the Data Model

The role definition is a property of the conversational session, not of the metric. Two sessions with the same conversation content but different assigned roles represent fundamentally different evaluation scenarios. Accordingly, `chatbot_role` is defined as a field on the `Dataset` object.

An alternative considered was to pass the role as a parameter at metric instantiation time, which would allow evaluating the same dataset against multiple roles. This was discarded on the grounds that a dataset already carries implicit contextual assumptions: the way a user interacts with a customer support agent differs from how they interact with a coding assistant, even if the raw conversation text were identical. Separating the role from the session would break this coupling and weaken the semantic integrity of the dataset object.

4.2 Format

The role definition is provided as a free-form natural language string. This string is passed directly to all scoring methods: embedded alongside the candidate response in deterministic path computations, and injected verbatim into the LLM judge prompt. No structured schema or explicit field decomposition is required.

5 Scoring Methodology

5.1 Scoring Strategy

The metric was originally designed with deterministic semantic scoring as its primary path, motivated by reproducibility and zero per-inference cost. Empirical evaluation (Section 8) showed that deterministic methods, while reliable at detecting semantic drift from a reference response, fail to detect tone violations and constructive failures — the hardest and most practically relevant violation types in a role-constrained deployment. This finding motivated adopting the LLM judge as the default scoring path.

The metric uses an LLM-as-a-judge as its default scoring path. The judge reads the role definition directly and is the only method capable of detecting tone violations and constructive failures — a conclusion supported empirically in Section 8.

When ground-truth reference responses are available and the user prefers a deterministic, reproducible alternative, the metric can be configured to use a semantic scoring function instead (see Section 10). Deterministic methods are faster and free of per-call cost, but detect semantic drift from the reference rather than role adherence per se.

5.2 LLM-as-a-Judge Path

Each assistant turn t_i is evaluated by an LLM judge that receives: (1) the role definition R; (2) the conversation history $T_{<i}$; and (3) the turn to evaluate t_i. The judge prompt implements the constructive adherence definition: the judge is asked not only whether the turn avoids scope violations, but whether it actively reflects the expected behaviour of the defined role.

Evaluation granularity. In the default `'turn'` mode the judge evaluates each turn independently with its prior context, producing n separate scores aligned with Equation 1. In `'conversation'` mode the entire conversation is submitted in a single call, reducing inference cost at the expense of per-turn granularity.

Output mode. The judge returns a **binary score** in $\{0, 1\}$ by default. Binary scoring is more consistent across calls; LLMs are generally not well-calibrated for numeric outputs. Continuous scoring in $[0, 1]$ is available as an option for deployments that prioritise granularity

over reproducibility. When `include_reason` is enabled, the judge additionally returns a natural-language justification for each per-turn score.

5.3 Deterministic Path

When a `ground_truth_assistant` response is available for each turn, six scoring functions are proposed: BERTScore, cosine similarity, k-nearest neighbours, and NLI compare individual turn pairs and are independent of corpus size. Mahalanobis Distance and Kullback–Leibler Divergence are distribution-based: they build a model of the role from the full corpus of ground-truth responses, and their signal quality varies with the number of reference examples available. A PCA-based variant of Mahalanobis Distance is additionally evaluated as a candidate solution to the curse-of-dimensionality problem and reported in Section 8. The per-turn score is continuous in $[0, 1]$ and is averaged across turns to produce the session score.

A note on scope. Deterministic methods compare candidate responses against a ground-truth reference and measure *semantic proximity*, not role adherence directly. A response can achieve high semantic similarity to the ground truth while violating the role's tone, register, or constructive requirements — violations that are semantically indistinguishable from adherent responses but contextually non-compliant. This limitation is characterised empirically in Section 8 and discussed in Section 9.

5.3.1 BERTScore

BERTScore [Zhang et al., 2020] computes semantic similarity between two texts by extracting contextual token embeddings from a pre-trained language model and computing precision, recall, and F1 over token-level cosine similarities. Unlike surface-level metrics, it captures paraphrastic equivalence. BERTScore F1 is computed as:

$$\text{BERTScore}_{F_1} = 2 \cdot \frac{P_{\text{BERT}} \cdot R_{\text{BERT}}}{P_{\text{BERT}} + R_{\text{BERT}}} \tag{2}$$

Role-adherent responses are not expected to reproduce the exact phrasing of a reference. BERTScore tolerates natural linguistic variation while remaining independent of a secondary judge model.

5.3.2 Semantic Cosine Similarity

Cosine similarity between sentence-level embeddings [Reimers and Gurevych, 2019] is a lighter-weight alternative. A sentence encoder maps each response to a dense vector; adherence is scored as the cosine of the angle between the candidate and reference vectors:

$$\text{sim}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \, \|\mathbf{b}\|} \tag{3}$$

Cosine similarity operates at sentence granularity rather than token granularity, making it faster. BERTScore and cosine similarity measure semantic proximity at different levels of granularity; Section 8 characterises their practical differences in group separation.

5.3.3 k-Nearest Neighbours

k-Nearest Neighbours (k-NN) generalises cosine similarity to multi-reference settings [Cover and Hart, 1967]. Given a reference set of k ground-truth embeddings for a turn, the adherence score is the cosine proximity to the nearest neighbour:

$$\text{kNN}(\mathbf{x}) = 1 - \min_{j \in \{1,\ldots,k\}} d_{\cos}(\mathbf{x}, \mathbf{g}_j) \tag{4}$$

where $d_{\cos}$ denotes cosine distance. When $k = 1$ and a single reference per turn is available, Equation 4 is mathematically equivalent to cosine similarity. The distinction materialises when multiple valid reference responses exist per turn: k-NN distances to the k nearest neighbours weight the score towards the most similar acceptable response, gaining robustness over a single-reference cosine comparison. We propose k-NN as the natural extension of cosine similarity for production deployments where multiple ground-truth responses per turn can be maintained.

5.3.4 Natural Language Inference

NLI models [Bowman et al., 2015] classify the relationship between two texts as *entailment*, *neutral*, or *contradiction*. Applied to role adherence, a candidate response that contradicts what a ground-truth role-adherent response asserts is a strong signal of non-adherence. NLI is proposed as a complementary signal for detecting explicit violations rather than measuring overall similarity. A *contradiction* label serves as a hard penalty regardless of similarity scores.

5.3.5 Mahalanobis Distance

Cosine similarity computes the angle between embedding vectors but treats all dimensions equally, regardless of how much the role actually varies along each dimension. A role that discusses diverse topics (high variance) but maintains strict tone constraints (low variance) will be mis-scored: large topic-dimensional shifts are penalised as heavily as small tone-dimensional shifts.

Mahalanobis Distance [Mahalanobis, 1936] corrects this by incorporating the empirical covariance structure of the role's embedding space. Given the set of ground-truth embeddings $\{\mathbf{g}_1, \ldots, \mathbf{g}_n\}$, let $\boldsymbol{\mu}$ denote their mean and $\boldsymbol{\Sigma}$ their covariance matrix. The Mahalanobis Distance of a candidate response embedding $\mathbf{x}$ from the role centroid is:

$$D_M(\mathbf{x}, \boldsymbol{\mu}) = \sqrt{(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})} \tag{5}$$

In practice $\boldsymbol{\Sigma}$ is approximated as a diagonal matrix (to avoid singularity in high-dimensional embedding spaces) and regularised as $\boldsymbol{\Sigma} + \lambda \mathbf{I}$. The distance is converted to an adherence score via score $= \exp(-\alpha \cdot D_M)$, where α is a scaling parameter set to 1 by default.

Sample sensitivity. Because $\boldsymbol{\Sigma}$ is estimated from the ground-truth corpus, signal quality varies with the number of available reference examples. Section 8 characterises this empirically across conditions.

5.3.6 Mahalanobis Distance with PCA Pre-Reduction

Mahalanobis Distance in high-dimensional embedding spaces is susceptible to the curse of dimensionality [Beyer et al., 1999]: as the number of dimensions grows, pairwise distances concentrate around a fixed value regardless of true class proximity, eroding discriminative power. Principal Component Analysis [Jolliffe, 2002] offers a candidate solution by projecting embeddings to a lower-dimensional subspace before applying Mahalanobis Distance.

We evaluate a pre-reduction from 384 to 32 dimensions (top-32 principal components of the ground-truth embedding matrix) followed by Mahalanobis Distance in the compressed space. This variant is evaluated empirically but **not proposed as a recommended scoring function**; Section 8 shows it produces inverted adherence scores and is reported as a negative finding.

5.3.7 Kullback–Leibler Divergence

For roles with explicit stylistic or vocabulary constraints, adherence can be framed as a distributional comparison. Let P denote the vocabulary distribution of the role estimated from

ground-truth responses, and let Q denote the vocabulary distribution of the candidate turn. The Kullback–Leibler Divergence [Kullback and Leibler, 1951] from Q to the reference P is:

$$D_{\mathrm{KL}}(P \,\|\, Q) = \sum_{w \in \mathcal{V}} P(w) \log \frac{P(w)}{Q(w)} \tag{6}$$

We use $D_{\mathrm{KL}}(P\|Q)$ — the divergence of the candidate from the role reference — because P is the distributional target. To handle zero probabilities, both distributions are smoothed with Laplace smoothing. The vocabulary $\mathcal{V}$ is restricted to the top-k most frequent tokens in the ground-truth corpus (default $k = 500$) to reduce sparsity. The adherence score is score $= \exp(-\beta \cdot D_{\mathrm{KL}})$, with $\beta = 1$ by default.

KL Divergence is proposed as a signal for vocabulary drift: cases where a response uses phrasing or terminology the role explicitly avoids. Section 8 characterises its empirical performance.

Sample sensitivity. Estimating P reliably requires a representative vocabulary distribution from the ground-truth corpus. Section 8 characterises how signal quality varies with corpus size.

5.3.8 Lexical Metrics: Considered and Discarded

ROUGE [Lin, 2004] and BLEU [Papineni et al., 2002] measure n-gram overlap between candidate and reference texts. Both were considered and discarded as primary metrics for this use case.

Role-adherent responses are not expected to reproduce the exact phrasing of a reference — they are expected to exhibit the same role-consistent behaviour with natural linguistic variation. N-gram metrics penalise paraphrases and reward surface-level copying, neither of which is desirable here. ROUGE and BLEU are noted as valid lightweight baselines for benchmarking purposes only.

6 Hypotheses

We formulate four hypotheses to guide the empirical evaluation in Sections 7–8.

H1 At small corpus sizes, KL Divergence does not produce a stable discriminative signal because the ground-truth vocabulary distribution is too sparse to be representative. We additionally predict, from the concentration of measure in high-dimensional sentence embedding spaces [Beyer et al., 1999], that Mahalanobis Distance will *degrade* as the corpus grows — a failure mode distinct from the sparse-distribution instability of KL Divergence.

H2 With a sufficient corpus, deterministic methods achieve agreement with the gold standard comparable to the LLM judge, while offering perfect reproducibility and zero per-inference cost.

H3 Mahalanobis Distance detects tone violations with higher precision than cosine similarity, by weighting low-variance stylistic dimensions more heavily than high-variance topical dimensions.

H4 KL Divergence provides effective signal for vocabulary drift violations that embedding-based similarity metrics fail to capture, because it operates directly on token-level distributions rather than aggregated dense representations.

H1 motivates evaluating distribution-based methods across corpus sizes rather than at a single fixed point. H2 motivated exploring deterministic methods as a cost-effective alternative to the LLM judge; Section 9 reports it refuted. H3 and H4 require per-violation-type evaluation and cannot be resolved from aggregate binary classification alone; this is addressed in Section 9.

7 Experimental Setup

7.1 Benchmark Dataset

We construct a controlled synthetic benchmark for a single role: a customer-facing support agent for a fintech application. The role is defined with four explicit constraints: (1) scope limited to account inquiries, transaction history, and card management; (2) a professional, empathetic, and concise tone; (3) no financial or investment advice under any circumstances; and (4) a constructive requirement to always offer a next step or resolution path rather than an open-ended refusal.

The benchmark consists of 150 assistant turns organised in 30 simulated conversations of 5 turns each. Labels are assigned by design: each turn is generated to exhibit one of four target behaviours — *adherent*, *scope violation*, *tone violation*, or *constructive failure* — and the label reflects the construction intent rather than post-hoc annotation. The distribution is approximately 40% adherent and 20% per violation type.

Each record contains four fields:

- **role**: the full role definition string.
- **assistant**: the generated response, which may violate the role.
- **ground_truth_assistant**: a role-adherent reference response.
- **label**: the designed ground-truth class.

7.2 Evaluation Conditions

Three conditions evaluate the same benchmark at increasing n:

- **Condition A** ($n = 15$): the first 15 turns.
- **Condition B** ($n = 50$): the first 50 turns.
- **Condition C** ($n = 150$): all 150 turns.

Conditions are cumulative: B includes all turns in A, and C includes all turns in B. This isolates the effect of dataset size on each scoring method while holding the role and violation distribution constant.

7.3 Scoring Methods

Table 1: Scoring methods evaluated, with their data requirements.

Method	Requires GT	Distribution-based
LLM judge (binary)	No	No
BERTScore F1	Yes	No
Cosine similarity	Yes	No
k-NN	Yes	No
NLI	Yes	No
Mahalanobis Distance	Yes	Yes
PCA+Mahalanobis (eval. only)	Yes	Yes
KL Divergence	Yes	Yes

7.4 Evaluation Protocol

Continuous methods (BERTScore, cosine similarity, k-NN, Mahalanobis Distance, PCA+Mahalanobis, and KL Divergence) require a threshold to produce binary predictions; any threshold is arbitrary in production. We report two threshold-agnostic statistics:

AUC-ROC [Hanley and McNeil, 1982] is the probability that the method assigns a higher score to a randomly selected adherent turn than to a randomly selected violation. AUC = 1.0 indicates perfect ranking; AUC = 0.5 is chance.

Separation is the difference between the mean score of adherent turns and the mean score of violation turns. It quantifies the absolute gap between group centroids and reflects how reliably a threshold can be set in practice. High AUC with low separation indicates perfect ranking but fragile threshold selection.

Binary classifiers (NLI and LLM Judge) produce hard predictions directly. We report three statistics. **F1 per class** [Manning et al., 2008] is the harmonic mean of precision and recall for each class:

$$F_1 = 2 \cdot \frac{P \cdot R}{P + R} \tag{7}$$

where P and R are precision and recall for the target class. **Macro F1** is the unweighted average of per-class F1 scores, treating each class equally regardless of frequency — a relevant choice here because the adherent and violation classes are not balanced. **Cohen's κ** [Cohen, 1960] is a chance-corrected agreement coefficient:

$$\kappa = \frac{p_o - p_e}{1 - p_e} \tag{8}$$

where p_o is the observed fraction of agreements between the classifier and the gold labels, and p_e is the fraction expected by chance under the marginal class distributions. $\kappa = 1.0$ indicates perfect agreement; $\kappa = 0$ indicates agreement at chance level; negative values indicate worse-than-chance performance.

For the LLM Judge, macro F1 and per-class F1 are additionally accompanied by 95% bootstrap confidence intervals (1 000 resamples, seed 42) to quantify estimation uncertainty at each sample size.

Mahalanobis Distance uses diagonal covariance estimation with ℓ_2 regularisation ($\lambda = 0.01$). KL Divergence uses Laplace-smoothed distributions restricted to the top-500 most frequent tokens in the ground-truth corpus. The LLM judge is run once per turn per condition at temperature 0.

8 Results

8.1 Continuous Scoring Methods

Table 2 reports mean score by class (adherent and violation), separation, and AUC-ROC for each continuous method across the three conditions. Figure 1 provides a visual comparison of per-class mean scores across methods and conditions.

Figure 2 shows the separation curve for Mahalanobis Distance, KL Divergence, and PCA+Mahalanobis over intermediate sample sizes from $n = 15$ to $n = 150$. Mahalanobis separation declines monotonically from +0.056 to +0.041. KL separation rises from near zero (+0.002) to +0.023. PCA+Mahalanobis separation remains negative at every n, converging from below towards zero without crossing it.

Table 2: Continuous scoring methods: mean score by class, separation, and AUC-ROC across evaluation conditions. Sep. = mean(adh.) − mean(viol.); higher is better. AUC: 1.0 = perfect ranking, 0.5 = chance. Bold indicates inverted signal.

Method	Cond.	Adh.	Viol.	Sep.	AUC
BERTScore	$n = 15$	0.941	0.864	+0.077	1.000
	$n = 50$	0.940	0.869	+0.071	1.000
	$n = 150$	0.941	0.868	+0.073	0.995
Cosine	$n = 15$	0.905	0.466	+0.439	1.000
	$n = 50$	0.902	0.520	+0.382	0.998
	$n = 150$	0.877	0.497	+0.380	0.994
k-NN	$n = 15$	0.905	0.496	+0.409	1.000
	$n = 50$	0.902	0.544	+0.358	0.998
	$n = 150$	0.877	0.546	+0.331	0.993
Mahalanobis	$n = 15$	0.404	0.348	+0.056	0.963
	$n = 50$	0.391	0.348	+0.043	0.872
	$n = 150$	0.391	0.350	+0.041	0.856
PCA+Maha	$n = 15$	0.283	0.419	**−0.136**	**0.000**
	$n = 50$	0.318	0.393	**−0.075**	**0.110**
	$n = 150$	0.343	0.383	**−0.040**	**0.219**
KL Div.	$n = 15$	0.860	0.858	+0.002	0.537
	$n = 50$	0.693	0.671	+0.022	0.870
	$n = 150$	0.517	0.494	+0.023	0.877

Table 3: Binary classifiers: per-class F1, macro F1, and Cohen's κ.

Method	Cond.	F1 adh.	F1 viol.	F1 macro	κ
NLI	$n = 15$	0.571	0.000	0.286	0.000
	$n = 50$	0.625	0.333	0.479	0.167
	$n = 150$	0.612	0.269	0.441	0.128
LLM Judge	$n = 15$	1.000	1.000	1.000	1.000
	$n = 50$	0.974	0.984	0.979	0.958
	$n = 150$	0.974	0.984	0.979	0.958

8.2 Binary Classifiers

8.3 LLM Judge: Bootstrap Confidence Intervals

The $n = 15$ result ($F_1 = 1.000$, $\kappa = 1.000$) reflects perfect classification on a small and easy subsample and should not be taken as the representative estimate. Table 4 reports macro F1, per-class F1, and Cohen's κ with 95% bootstrap confidence intervals (1 000 resamples, seed 42). The canonical result is $n = 150$: $F_1^{\text{macro}} = 0.979$, $\kappa = 0.959$, CI $[0.952, 1.000]$.

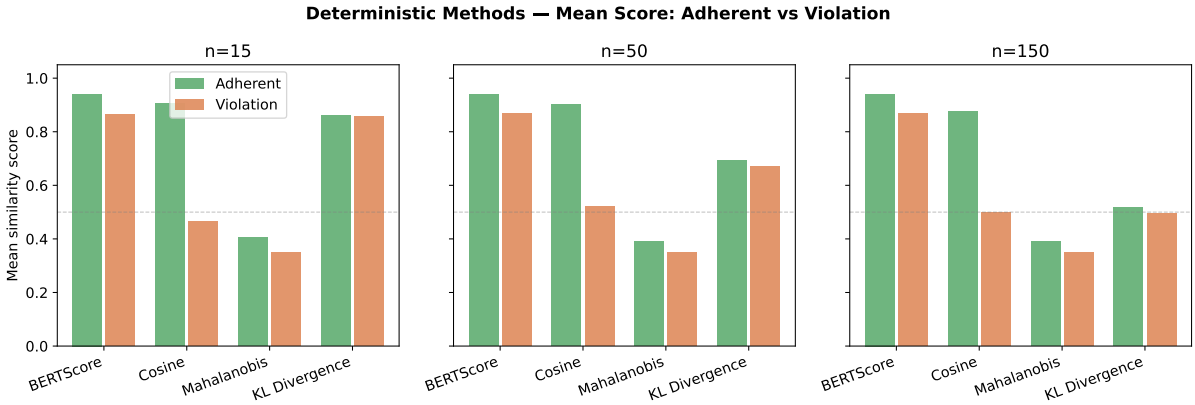

Figure 1: Mean similarity score per class (adherent vs. violation) for BERTScore, Cosine similarity, Mahalanobis Distance, and KL Divergence across the three evaluation conditions.

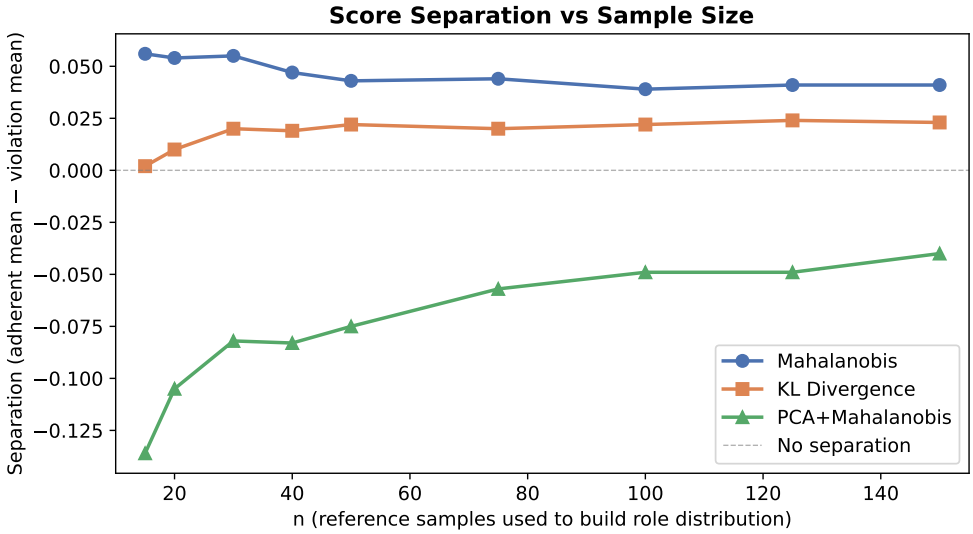

Figure 2: Score separation (adherent mean − violation mean) as a function of sample size for Mahalanobis Distance, KL Divergence, and PCA+Mahalanobis. Positive values indicate the method ranks adherent responses higher on average; negative values indicate signal inversion.

9 Discussion

9.1 The Semantics Trap: Why AUC ≈ 1 Is Not a Victory

BERTScore, Cosine, and k-NN achieve AUC ≈ 1.0 across all conditions. Read naively, this looks like a definitive result. It is not. The benchmark was constructed so that adherent responses are paraphrases of the ground truth (semantically close by design) and violations are responses that mishandle or ignore the request (semantically distinct by design). Any method that measures semantic proximity to a per-turn ground-truth reference separates this benchmark perfectly because separability was built into the construction, not discovered by the method.

The per-class mean scores make the mechanism explicit: Cosine similarity sits at 0.877 for adherent responses and 0.497 for violations at $n = 150$ — the groups are nearly half a unit apart because adherent examples were generated as paraphrases of the ground truth and violations were not. In production, the hardest violations — tone deviations and constructive failures — are precisely the ones where semantic similarity to the ground truth is highest: a response may say the right things in the wrong register, or omit the required resolution path while remaining

Table 4: LLM Judge (`llama-3.3-70b-versatile`): macro F1, per-class F1, and Cohen's κ with 95% bootstrap confidence intervals (1 000 resamples, seed 42).

Cond.	F1 macro	95% CI	F1 adh.	95% CI	F1 viol.	95% CI
$n = 15$	1.000	$[1.000, 1.000]$	1.000	$[1.000, 1.000]$	1.000	$[1.000, 1.000]$
$n = 50$	0.979	$[0.934, 1.000]$	0.975	$[0.914, 1.000]$	0.984	$[0.947, 1.000]$
$n = 150$	0.979	$[0.952, 1.000]$	0.975	$[0.942, 1.000]$	0.984	$[0.963, 1.000]$

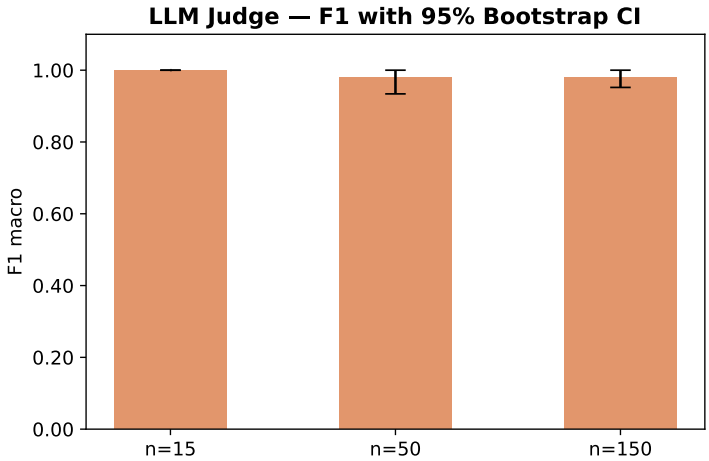

Figure 3: LLM Judge macro F1 with 95% bootstrap confidence intervals across evaluation conditions.

topically on-role. Cosine similarity would assign these a high score and miss the violation entirely.

The correct interpretation of deterministic methods with ground truth is that they are *semantic drift detectors*, not role adherence evaluators. They are useful for monitoring whether response content deviates from an expected reference, but they cannot substitute for a method that reads and understands the role definition.

9.2 Cosine vs. BERTScore: Practical Separation Matters

Of the methods that reach AUC ≈ 1.0, Cosine similarity achieves a separation of $+0.380$ at $n = 150$ while BERTScore achieves only $+0.073$. Although BERTScore operates at finer granularity (token-level contextual embeddings versus sentence-level pooling), this additional resolution does not translate to better group separation in this setup: the two groups are sufficiently similar at the token level that BERTScore barely distinguishes them in absolute terms, even though its ranking remains nearly perfect.

For a deployment requiring threshold-based binary classification, BERTScore's small separation makes threshold selection fragile: a threshold at 0.90 and one at 0.91 produce meaningfully different decisions. Cosine similarity, with its large separation, is more robust to threshold choice.

k-NN is mathematically equivalent to Cosine in the single-reference-per-turn setup used here ($k = 1$). Its advantage materialises when multiple valid ground-truth responses exist per turn: the nearest-neighbour distance weights the score towards the most similar acceptable response, gaining robustness that a single-reference cosine comparison cannot provide.

9.3 Mahalanobis Distance: Curse of Dimensionality Confirmed

Mahalanobis Distance shows a monotonic degradation pattern: AUC is highest at $n = 15$ (0.963) and declines to 0.856 at $n = 150$. This is the opposite of what standard statistical intuition would predict — more data should produce better distributional estimates, not worse ones.

The explanation is the concentration of measure in high-dimensional spaces [Beyer et al., 1999]. As n grows, the estimated covariance of the 384-dimensional ground-truth embeddings converges towards the true distribution, but in 384D all pairwise distances concentrate around a fixed value regardless of class membership. The diagonal regularisation ($\lambda = 0.01$) partially mitigates this at small n by preventing covariance collapse; at larger n the regularisation becomes relatively negligible and the concentration effect dominates. This confirms the prediction embedded in H1: Mahalanobis Distance is not viable for production deployments at standard embedding dimensions.

9.4 PCA+Mahalanobis: A Structural Failure

The PCA+Mahalanobis variant was motivated as a dimensionality reduction solution to the concentration problem. The result is unambiguous: the method inverts the adherence signal at every condition (sep $= -0.136$ at $n = 15$, never crossing zero). AUC $= 0.000$ at $n = 15$ means the method assigns higher scores to violations than to adherent responses with perfect consistency.

The mechanism is structural. PCA decomposes the variance of the ground-truth embedding matrix: the 32 retained principal components capture the directions of maximum variance within the reference set. This transformation has no relationship to the adherence-violation geometry — it captures what varies most *among adherent examples*, not what separates adherent from violation. In the compressed space, violation embeddings happen to project near the centroid of the transformed reference distribution while adherent examples, spread across the principal components, project further away.

More data partially corrects the PCA estimate (components become less dominated by individual examples), which is why separation improves from -0.136 to -0.040 as n grows. But the structural distortion never reverses. PCA+Mahalanobis is **not recommended** for any deployment scenario. H2 is **refuted**: no deterministic method approaches the LLM Judge at large n, and the methods that appear to do so via AUC are benefiting from the benchmark artefact described above.

9.5 KL Divergence: The Wrong Level of Abstraction

KL Divergence operates on vocabulary distributions estimated from the ground-truth corpus. Its near-random signal at $n = 15$ (AUC $= 0.537$) is expected: with 15 turns the top-500 vocabulary distribution is too sparse to be representative. As n grows the distribution stabilises and AUC converges to 0.877 at $n = 150$.

The ceiling is not a sample-size problem — it is a conceptual one. The role violations in the benchmark (scope, tone, constructive failures) do not produce distinctive vocabulary patterns. An agent that answers a question in the wrong tone uses largely the same vocabulary as one that answers correctly. KL Divergence is blind to this distinction because it operates at the lexical level, not the semantic or pragmatic level. H4 cannot be confirmed from aggregate binary classification; the convergence to AUC $= 0.877$ suggests KL may serve as a rough vocabulary drift monitor but not as a primary adherence detector.

9.6 NLI: Insufficient for Role Adherence

NLI achieves $F_1^{\text{macro}} = 0.441$ and $\kappa = 0.128$ at $n = 150$ — better than chance but far below useful production thresholds. The model correctly flags some explicit scope violations (responses that contradict ground-truth assertions) but misses tone violations and constructive failures. These

failure types do not produce logical contradictions with the ground-truth response: an agent that fails to offer a resolution path omits something, it does not contradict anything. NLI's contradiction signal is a necessary but insufficient condition for role violation. Performance is also unstable across n: κ drops from 0.167 at $n = 50$ to 0.128 at $n = 150$, suggesting the threshold used for binary classification is not stable.

9.7 LLM Judge: The Only Role-Aware Method

The LLM Judge achieves $F_1^{\text{macro}} = 0.979$ at $n \geq 50$ with bootstrap CI $[0.952, 1.000]$ at $n = 150$, $\kappa = 0.959$ $[0.904, 1.000]$. It is the only evaluated method that reads the role definition directly and evaluates the response against it. All other methods operate on signals that are at best correlated with adherence; the judge assesses adherence by definition and can detect tone violations and constructive failures that are invisible to semantic similarity methods.

The practical costs are non-determinism (mitigated by fixing temperature to 0) and per-call inference cost, both well-documented properties of LLM evaluation [Zheng et al., 2023].

Judge model quality is a design variable. The $F_1 = 0.979$ result was obtained with `llama-3.3-70b-versatile`. A smaller model will produce substantially different results. Reporting LLM Judge performance without specifying the judge model is methodologically equivalent to reporting a deterministic metric without specifying its embedding model: a necessary parameter omitted. H3 cannot be directly tested with aggregate binary classification; per-violation-type evaluation is identified as a direction for future work.

10 Configuration Parameters

The metric exposes the following parameters.

`scoring_method` (default: `'llm_judge'`) — selects the scoring function. `'llm_judge'` (default) reads the role definition directly and is the only option capable of detecting tone violations and constructive failures. When ground-truth responses are available, deterministic alternatives can be selected: `'cosine'` and `'knn'` are recommended for cost-sensitive monitoring, with the explicit caveat that they detect semantic drift from the reference rather than role adherence per se; `'bertscore'` and `'nli'` are also available. See Section 8 for empirical characterisation of each method.

`output_mode` (default: `'binary'`) — controls the LLM judge output format. `'binary'` returns a hard $\{0, 1\}$ score per turn; `'continuous'` returns a float in $[0, 1]$. Binary scoring is more consistent across calls. Continuous scoring is more expressive but requires deliberate threshold selection. Applies only when `scoring_method='llm_judge'`.

`evaluation_granularity` (default: `'turn'`) — controls whether the judge evaluates each assistant turn with its prior context (`'turn'`, one API call per turn, aligned with Equation 1) or the full conversation in a single call (`'conversation'`, lower cost, no per-turn breakdown). Applies only when `scoring_method='llm_judge'`.

`strict_mode` (default: `False`) — when `True`, the session passes only if every turn is adherent (session score = 1.0). Intended for high-criticality deployments where partial adherence is not acceptable.

`threshold` (default: 0.5) — the minimum session score required to mark the conversation as adherent. Ignored when `strict_mode=True`. In binary output mode a threshold of 0.5 requires that at least half of all turns are adherent; most relevant in `output_mode='continuous'`.

`model` — the LLM used as judge. Follows the `Guardian` abstraction in the Gaussia framework, allowing any supported provider to be substituted. As shown in Section 9, judge model quality significantly affects detection performance and should be treated as a first-class design parameter.

`include_reason` (default: `False`) — when enabled, the judge returns a natural-language justification alongside each per-turn score. Applies only when `scoring_method='llm_judge'`.

11 Implementation Considerations

The metric extends the `Gaussia` base class. The `batch()` method iterates over each assistant turn in a session, dispatches to the scoring function selected via `scoring_method` (default: `'llm_judge'`), and appends a per-turn score. The session score is the mean of per-turn scores (Equation 1).

The `chatbot_role` field is read from the `Dataset` object and passed as a string to all scoring methods and, when the judge path is active, injected verbatim into the judge prompt.

Statistical mode compatibility (Frequentist and Bayesian) follows the existing Gaussia pattern. No changes to the statistical layer are required.

12 Future Work

Two directions are identified for future development.

Multimodal settings. As conversational AI systems process images, audio, and video alongside text, role adherence evaluation must generalise beyond text-only responses. The LLM judge path extends naturally by requiring a vision-capable model; the deterministic path would require modality-appropriate similarity functions, such as CLIP-based similarity [Radford et al., 2021] for image content.

Re-evaluation on a larger and more diverse benchmark. The current benchmark is synthetic and scoped to a single domain. As discussed in Section 9, the near-perfect AUC of semantic methods is a structural property of this benchmark's construction rather than evidence of genuine role understanding. A multi-domain benchmark with naturally occurring violations — including tone deviations that are semantically close to adherent responses — would provide a more demanding test for all methods and may yield a different relative ordering.

13 Conclusion

We have proposed Role Adherence, a metric for the Gaussia evaluation framework that measures how consistently an AI assistant maintains its assigned role across a multi-turn conversation. The metric adopts a constructive definition of adherence, evaluates each turn in full conversational context, and uses an LLM-as-a-judge as its default scoring path. Deterministic semantic scoring is available as a reproducible alternative when ground-truth reference responses are present.

Empirical evaluation on a controlled synthetic benchmark surfaces a distinction with practical consequences: deterministic methods (BERTScore, Cosine, k-NN) measure *semantic proximity* to a reference response, not role adherence per se. Their near-perfect AUC on the benchmark is a structural artefact of how the benchmark was constructed, not evidence of role understanding. A response that violates the role in tone or fails a constructive requirement while remaining semantically close to the expected answer is invisible to these methods. This finding motivates using deterministic methods only as lightweight semantic drift monitors, not as primary adherence evaluators.

Among methods that model the reference distribution rather than individual turn pairs, Mahalanobis Distance in high-dimensional embedding space degrades monotonically with sample size (AUC $0.963 \rightarrow 0.856$), confirming the curse of dimensionality. A PCA-based pre-reduction inverts the adherence signal entirely and is not recommended. KL Divergence converges to a weak ceiling (AUC 0.877) because vocabulary overlap is a poor proxy for adherence. The LLM Judge ($F_1 = 0.979$, $\kappa = 0.959$, CI $[0.952, 1.000]$ at $n = 150$) is the only method that reliably

detects all violation types, at the cost of per-call inference and dependence on judge model quality — a variable that must be treated as a first-class design parameter.

Primary directions for future work are: (1) per-violation-type evaluation to test whether any deterministic method retains diagnostic value for specific violation categories; (2) multi-domain benchmarking beyond the single fintech support agent domain used here; and (3) cost-accuracy analysis of judge models to characterise the tradeoff between inference cost and detection quality.

References

Kevin Beyer, Jonathan Goldstein, Raghu Ramakrishnan, and Uri Shaft. When is "nearest neighbor" meaningful? In *Proceedings of the 7th International Conference on Database Theory*, volume 1540 of *Lecture Notes in Computer Science*, pages 217–235, 1999.

Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. A large annotated corpus for natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 632–642, 2015.

Jacob Cohen. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46, 1960.

Thomas M. Cover and Peter E. Hart. Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1):21–27, 1967.

Jan Deriu, Alvaro Rodrigo, Arantxa Otegi, Guillermo Echegoyen, Sophie Rosset, Eneko Agirre, and Mark Cieliebak. Survey on evaluation methods for dialogue systems. *Artificial Intelligence Review*, 54(1):755–810, 2021.

James A. Hanley and Barbara J. McNeil. The meaning and use of the area under a receiver operating characteristic (ROC) curve. *Radiology*, 143(1):29–36, 1982.

Ian T. Jolliffe. *Principal Component Analysis*. Springer, New York, 2 edition, 2002.

Solomon Kullback and Richard A. Leibler. On information and sufficiency. *The Annals of Mathematical Statistics*, 22(1):79–86, 1951.

Chin-Yew Lin. ROUGE: A package for automatic evaluation of summaries. In *Text Summarization Branches Out: Proceedings of the ACL-04 Workshop*, pages 74–81, 2004.

Prasanta Chandra Mahalanobis. On the generalised distance in statistics. *Proceedings of the National Institute of Sciences of India*, 2(1):49–55, 1936.

Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. Cambridge University Press, Cambridge, 2008.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, 2002.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, pages 8748–8763, 2021.

Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 3982–3992, 2019.

Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. BERTScore: Evaluating text generation with BERT. In *International Conference on Learning Representations*, 2020.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, volume 36, 2023.